

# Sentiment Analysis

## 1.1) Load & Combine Data

```
In [1]: import pandas as pd # for handling dataframes
import re # for regular expressions
import numpy as np # for numerical operations
from collections import Counter # for counting words/n-grams
from sklearn.feature_extraction.text import ENGLISH_STOP_WORDS # stopwords list

# --- Load Data from text files ---
# Each dataset has sentences and labels (0=negative, 1=positive)

df_yelp = pd.read_csv(
    './sentiment labelled sentences/yelp_labelled.txt', # path to Yelp dataset
    names=['sentence', 'label'], # column names
    sep='\t', # tab-separated values
    encoding='utf-8'
)

df_amazon = pd.read_csv(
    './sentiment labelled sentences/amazon_cells_labelled.txt', # Amazon reviews
    names=['sentence', 'label'],
    sep='\t',
    encoding='utf-8'
)

df_imdb = pd.read_csv(
    './sentiment labelled sentences/imdb_labelled.txt', # IMDB reviews
    names=['sentence', 'label'],
    sep='\t',
    encoding='utf-8'
)

# Print dataset shapes to see how many rows/examples each has
print("Yelp:", df_yelp.shape)
print("Amazon:", df_amazon.shape)
print("IMDB:", df_imdb.shape)

# --- Combine all datasets into one dataframe ---
df_all = pd.concat([df_yelp, df_amazon, df_imdb], ignore_index=True)
print("Combined:", df_all.shape) # check total number of examples

# Keep a copy of combined dataset for further processing
df = df_all.copy()
```

```
Yelp: (1000, 2)
Amazon: (1000, 2)
IMDB: (748, 2)
Combined: (2748, 2)
```

## 1.2) Explore Data

```
In [2]: # --- Tokenizer ---
# This regex keeps words and contractions like "don't"
_token_re = re.compile(r"[A-Za-z]+(?:'[A-Za-z]+)?")

def tok(s):
    """Convert a sentence to a list of lowercase tokens/words."""
    if not isinstance(s, str):
        return []
    return _token_re.findall(s.lower())

# --- Label distribution per dataset ---
for name, d in [('Yelp', df_yelp), ('Amazon', df_amazon), ('IMDB', df_imdb)]:
    percent = (d['label']
                .value_counts(normalize=True) # get percentages
                .sort_index() * 100).round(2)
    print(f"\n{name} label distribution (%):")
    print(percent)

# --- Overview of combined dataset ---
print("\nDataset shape:", df.shape) # number of rows and columns
print("Missing values:", df.isna().sum().to_dict()) # check for NaNs

# --- Global Label distribution ---
counts = df['label'].value_counts().sort_index() # counts of 0 and 1
pct = (df['label'].value_counts(normalize=True).sort_index() * 100).round(2) # per
print("\nGlobal label distribution:")
print(pd.DataFrame({'count': counts, 'percent': pct}))

# --- Basic text statistics ---
def stats(text):
    """Compute simple statistics for a sentence."""
    s = text if isinstance(text, str) else "" # handle non-string entries
    t = tok(s) # tokenize
    n = len(t) # number of tokens
    return pd.Series({
        'char_len': len(s), # number of characters
        'word_len': n, # number of words
        'avg_word_len': (sum(len(x) for x in t) / n) if n else 0.0, # average word
        'stop_ratio': sum(x in ENGLISH_STOP_WORDS for x in t) / (n or 1), # fracti
        'ttr': len(set(t)) / (n or 1), # type-token ratio (vocabulary richness)
    })

# Apply stats to all sentences
df_e = df['sentence'].apply(stats)
# Combine label column with computed stats
df_e = pd.concat([df[['label']], df_e], axis=1)

print("\nOverall descriptors:")
```

```
print(df_e.describe().round(3).T) # summary statistics

print("\nMean descriptors per label:")
print(df_e.groupby('label').mean(numeric_only=True).round(3)) # average per sentiment

# --- N-Gram Extraction ---
def top_ngrams(texts, n=1, k=10, drop_stop=True):
    """
    Extract top-k n-grams from a list of sentences.

    Args:
        texts: list of sentences
        n: size of n-gram (1=unigram, 2=bigram, etc.)
        k: top-k n-grams to return
        drop_stop: whether to remove stopwords
    Returns:
        DataFrame with top-k n-grams and counts
    """
    c = Counter()
    for s in texts:
        tokens = tok(s)
        if drop_stop:
            tokens = [w for w in tokens if w not in ENGLISH_STOP_WORDS]
        grams = tokens if n == 1 else [
            " ".join(tokens[i:i + n]) for i in range(len(tokens) - n + 1)
        ]
        c.update(grams)
    df_out = pd.DataFrame(c.most_common(k), columns=[f'{n}-gram', 'count'])
    return df_out

print("\nTop Unigrams (all):")
print(top_ngrams(df['sentence'], n=1, k=10)) # top 10 single words

print("\nTop Bigrams (all):")
print(top_ngrams(df['sentence'], n=2, k=10)) # top 10 two-word combinations

# Top n-grams per Label
for lbl in sorted(df['label'].unique()):
    print(f"\nTop Unigrams (Label={lbl}):")
    print(top_ngrams(df[df['label'] == lbl]['sentence'], n=1, k=10))

    print(f"\nTop Bigrams (Label={lbl}):")
    print(top_ngrams(df[df['label'] == lbl]['sentence'], n=2, k=10))

# --- Vocabulary Size ---
def vocab_size(texts):
    """Compute the number of unique words in a list of sentences, excluding stopwords"""
    v = set()
    for s in texts:
        v.update([w for w in tok(s) if w not in ENGLISH_STOP_WORDS])
    return len(v)
```

```
print("\nTotal vocabulary size (no stopwords):", vocab_size(df['sentence']))  
print("\nVocabulary size per label:")  
print(df.groupby('label')['sentence'].apply(vocab_size).to_frame('vocab_size'))
```

Yelp label distribution (%):  
label  
0 50.0  
1 50.0  
Name: proportion, dtype: float64

Amazon label distribution (%):  
label  
0 50.0  
1 50.0  
Name: proportion, dtype: float64

IMDB label distribution (%):  
label  
0 48.4  
1 51.6  
Name: proportion, dtype: float64

Dataset shape: (2748, 2)  
Missing values: {'sentence': 0, 'label': 0}

Global label distribution:  
count percent  
label  
0 1362 49.56  
1 1386 50.44

Overall descriptors:

	count	mean	std	min	25%	50%	75%	max
label	2748.0	0.504	0.500	0.0	0.000	1.000	1.000	1.0
char_len	2748.0	71.620	204.405	7.0	32.000	55.000	87.000	8041.0
word_len	2748.0	12.896	33.489	0.0	6.000	10.000	16.000	1297.0
avg_word_len	2748.0	4.498	0.974	0.0	3.923	4.364	4.909	13.0
stop_ratio	2748.0	0.494	0.181	0.0	0.417	0.500	0.609	1.0
ttr	2748.0	0.956	0.073	0.0	0.917	1.000	1.000	1.0

Mean descriptors per label:

label	char_len	word_len	avg_word_len	stop_ratio	ttr
0	74.437	13.467	4.462	0.515	0.955
1	68.852	12.334	4.533	0.474	0.956

Top Unigrams (all):  
1-gram count  
0 good 231  
1 great 210  
2 movie 181  
3 phone 165  
4 film 160  
5 food 126  
6 like 125  
7 just 119  
8 place 114  
9 it's 114

Top Bigrams (all):

	2-gram	count
0	waste time	17
1	works great	17
2	customer service	15
3	sound quality	14
4	waste money	12
5	don't waste	11
6	highly recommend	11
7	battery life	11
8	don't think	10
9	i've seen	10

Top Unigrams (Label=0):

	1-gram	count
0	bad	96
1	movie	95
2	phone	78
3	film	71
4	just	69
5	like	67
6	food	66
7	time	62
8	don't	60
9	good	57

Top Bigrams (Label=0):

	2-gram	count
0	waste time	16
1	customer service	12
2	waste money	12
3	don't waste	11
4	don't think	8
5	i've seen	8
6	it's just	6
7	doesn't work	6
8	does work	6
9	don't buy	6

Top Unigrams (Label=1):

	1-gram	count
0	great	193
1	good	174
2	film	89
3	phone	87
4	movie	86
5	food	60
6	really	60
7	it's	60
8	best	59
9	place	58

Top Bigrams (Label=1):

	2-gram	count
0	works great	17
1	highly recommend	11
2	sound quality	10

```
3      great phone      9
4      great food       8
5      really good     8
6      great service    7
7      pretty good      7
8      food good        7
9      good quality      7
```

Total vocabulary size (no stopwords): 4903

Vocabulary size per label:

label	vocab_size
0	3111
1	3079

## 2.1) Text Cleaning (Improved)

```
In [3]: import re # regular expressions for text processing
import html # for handling HTML entities

# --- Optional: Lemmatization using NLTK ---
use_lemmatize = True
try:
    import nltk
    from nltk.stem import WordNetLemmatizer

    # Check if NLTK data is downloaded
    nltk.data.find('corpora/wordnet')
    nltk.data.find('corpora/omw-1.4')
    _lemmatizer = WordNetLemmatizer() # create lemmatizer
except Exception:
    # fallback to stemming if NLTK or wordnet not available
    use_lemmatize = False

from nltk.stem import PorterStemmer # fallback stemmer

_stemmer = PorterStemmer() # create stemmer instance

# --- Common English contractions ---
# This dictionary maps contractions to their expanded form
# e.g., "don't" -> "do not"
CONTRACTIONS = {
    "ain't": "am not", "aren't": "are not", "can't": "cannot", "can't've": "cannot",
    "could've": "could have", "couldn't": "could not", "didn't": "did not", "doesn't": "does not",
    "don't": "do not", "hadn't": "had not", "hasn't": "has not", "haven't": "have not",
    "he's": "he is", "she's": "she is", "it's": "it is", "i'm": "i am", "i've": "i have",
    "i'd": "i would", "i'll": "i will", "isn't": "is not", "let's": "let us", "migh't": "might",
    "mustn't": "must not", "shan't": "shall not", "shouldn't": "should not",
    "that's": "that is", "there's": "there is", "they're": "they are", "they've": "they have",
    "we're": "we are", "we've": "we have", "we'll": "we will", "weren't": "were not",
    "what's": "what is", "who's": "who is", "won't": "will not", "wouldn't": "would not",
    "you'd": "you would", "you'll": "you will", "you're": "you are", "y'all": "you all"
}
```

```
# --- Regex patterns for cleaning ---
_url_re = re.compile(r"https?://\S+|www\.\S+", re.IGNORECASE) # match URLs
_html_tag_re = re.compile(r"<[^>]+>") # match HTML tags
_num_re = re.compile(r"\b\d+(?:[.,]\d+)?\b") # match numbers
_punct_re = re.compile(r"[\\.,!?:;\\-\\(\\))\\[\\]]\\{\\}\\\"`~\\\\@#\\$%\\^&\\*\\+=>\\]+") #
_token_re = re.compile(r"[A-Za-z]+(?:'[A-Za-z]+)?") # keep words and contractions

# --- Core functions ---

def expand_contractions(text: str) -> str:
    """Expand common English contractions to full words."""
    if not isinstance(text, str):
        return "" # return empty string if input is not text

    def repl(match):
        return CONTRACTIONS.get(match.group(0).lower(), match.group(0)) # replace

    # Create regex pattern from all contractions keys
    pattern = r"\b(" + "|".join(map(re.escape, CONTRACTIONS.keys())) + r")\b"
    return re.sub(pattern, repl, text, flags=re.IGNORECASE) # apply replacement

def basic_clean(text: str) -> str:
    """Clean text: remove URLs, HTML tags, numbers, punctuation, lowercase."""
    if not isinstance(text, str):
        return ""
    s = html.unescape(text) # convert HTML entities to normal text
    s = _url_re.sub(" ", s) # remove URLs
    s = _html_tag_re.sub(" ", s) # remove HTML tags
    s = expand_contractions(s) # expand contractions
    s = _num_re.sub(" ", s) # remove numbers
    s = _punct_re.sub(" ", s) # remove punctuation
    s = s.lower() # convert to lowercase
    s = re.sub(r"\s+", " ", s).strip() # remove extra spaces
    return s

def tokenize(text: str):
    """Split text into list of words (tokens)."""
    return _token_re.findall(text)

def normalize_tokens(tokens):
    """Convert tokens to base form using lemmatization or stemming."""
    out = []
    for t in tokens:
        if use_lemmatize:
            try:
                out.append(_lemmatizer.lemmatize(t)) # Lemmatize word
            except Exception:
                out.append(_stemmer.stem(t)) # fallback to stemming
        else:
            out.append(_stemmer.stem(t)) # stem if lemmatizer not available
    return out
```

```

def clean_text_pipeline(text: str, remove_stopwords=False):
    """
    Full text preprocessing pipeline.

    Args:
        text: raw sentence
        remove_stopwords: True to remove common English stopwords
    Returns:
        clean string, list of normalized tokens
    """
    clean_str = basic_clean(text) # clean text
    tokens = tokenize(clean_str) # tokenize
    if remove_stopwords:
        tokens = [t for t in tokens if t not in ENGLISH_STOP_WORDS] # remove stopw
    tokens = normalize_tokens(tokens) # lemmatize or stem
    return clean_str, tokens

# --- Apply cleaning to the dataset ---
# For deep Learning, keep stopwords (important for negations like "not good")
df['clean_text'], df['tokens'] = zip(*df['sentence'].map(lambda x: clean_text_pipeline(x)))

# --- Compute sentence length after cleaning ---
df['len_tokens'] = df['tokens'].apply(len) # number of tokens per sentence

# --- Quick check of cleaned data ---
print(df[['sentence', 'clean_text', 'label', 'len_tokens']].head()) # show first 5

```

	sentence \
0	Wow... Loved this place.
1	Crust is not good.
2	Not tasty and the texture was just nasty.
3	Stopped by during the late May bank holiday of...
4	The selection on the menu was great and so wer...

	clean_text	label	len_tokens
0	wow loved this place	1	4
1	crust is not good	0	4
2	not tasty and the texture was just nasty	0	8
3	stopped by during the late may bank holiday of...	1	15
4	the selection on the menu was great and so wer...	1	12

## 2.2) Exploratory Data Analysis (EDA) (Improved)

```

In [4]: import matplotlib.pyplot as plt # for plotting
import seaborn as sns # for nicer statistical plots
from wordcloud import WordCloud # for creating word clouds

# Set seaborn style for all plots (white background + muted colors + slightly Large
sns.set(style="whitegrid", palette="muted", font_scale=1.1)

# --- Label distribution by source ---
# Create a DataFrame with proportion of each label per dataset
label_df = pd.DataFrame({

```

```

'Yelp': df_yelp['label'].value_counts(normalize=True).sort_index(),
'Amazon': df_amazon['label'].value_counts(normalize=True).sort_index(),
'IMDB': df_imdb['label'].value_counts(normalize=True).sort_index()
}).T # transpose so that datasets are rows
label_df.columns = ['Negative (0)', 'Positive (1)'] # rename columns

# Plot grouped bar chart
label_df.plot(kind='bar', figsize=(8,5), color=['#6baed6', '#fc9272'])
plt.ylabel('Proportion')
plt.title('Label Distribution by Dataset Source')
plt.xticks(rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7) # horizontal grid lines
plt.tight_layout()
plt.show()

# --- Sentence Length distribution by Label ---
# Violin plot shows distribution and density of sentence lengths
plt.figure(figsize=(8,5))
sns.violinplot(x='label', y='len_tokens', data=df, palette=['#6baed6', '#fc9272'])
plt.title('Sentence Length Distribution by Label')
plt.xlabel('Label')
plt.ylabel('Number of Tokens')
plt.tight_layout()
plt.show()

# --- Top N-grams per Label (horizontal barplots) ---
def plot_top_ngrams_bar(texts, label_name, n=1, top_k=10, drop_stop=True):
    """Plot top-k unigrams or bigrams for a given label."""
    c = Counter() # count occurrences
    for s in texts:
        tokens = tok(s) # tokenize sentence
        if drop_stop:
            tokens = [w for w in tokens if w not in ENGLISH_STOP_WORDS] # remove stop words
        grams = tokens if n == 1 else [
            " ".join(tokens[i:i+n]) for i in range(len(tokens)-n+1) # create n-grams
        ]
        c.update(grams) # update counts
    top = c.most_common(top_k) # get top k n-grams
    if not top:
        return
    words, counts = zip(*top)
    plt.figure(figsize=(6,4))
    plt.barh(words[::-1], counts[::-1], color='#6baed6') # horizontal bar plot
    plt.title(f'Top {top_k} {"Unigrams" if n==1 else "Bigrams"} - Label={label_name}')
    plt.xlabel('Frequency')
    plt.tight_layout()
    plt.show()

# Plot top unigrams and bigrams for each label
for lbl in sorted(df['label'].unique()):
    plot_top_ngrams_bar(df[df['label']==lbl]['sentence'], label_name=lbl, n=1, top_k=10)
    plot_top_ngrams_bar(df[df['label']==lbl]['sentence'], label_name=lbl, n=2, top_k=10)

# --- Vocabulary overlap (Venn diagram) ---
try:
    from matplotlib_venn import venn2

```

```
# Get unique tokens per label (exclude stopwords)
tokens_neg = set([t for toks in df[df['label']==0]['tokens'] for t in toks if t
tokens_pos = set([t for toks in df[df['label']==1]['tokens'] for t in toks if t

plt.figure(figsize=(6,6))
venn2([tokens_neg, tokens_pos], set_labels=('Negative (0)', 'Positive (1)'),
      set_colors=('#fc9272', '#6baed6'), alpha=0.7)
plt.title('Vocabulary Overlap Between Labels')
plt.show()
except Exception as e:
    print("Venn diagram could not be plotted:", e)

# --- Label distribution (overall) ---
label_counts = df['label'].value_counts().sort_index()
plt.figure(figsize=(6,4))
sns.barplot(x=label_counts.index, y=label_counts.values, palette=['#6baed6', '#fc9272'])
plt.title('Label Distribution', fontsize=14)
plt.xlabel('Label', fontsize=12)
plt.ylabel('Count', fontsize=12)
# Add counts on top of bars for clarity
for i, v in enumerate(label_counts.values):
    plt.text(i, v + max(label_counts.values)*0.01, str(v), ha='center', fontsize=10)
plt.tight_layout()
plt.show()

# --- Word frequency plots ---
def plot_top_words(token_series, title, top_k=20, color="#74c476"):
    """Plot top-K words from tokenized series."""
    freq = Counter([t for tokens in token_series for t in tokens]) # count all tokens
    top = freq.most_common(top_k)
    if top:
        words, counts = zip(*top)
        plt.figure(figsize=(12,5))
        sns.barplot(x=list(words), y=list(counts), palette=[color]*len(words))
        plt.xticks(rotation=45, ha='right')
        plt.title(title, fontsize=14)
        plt.ylabel('Frequency', fontsize=12)
        plt.xlabel('')
        plt.tight_layout()
        plt.show()
    else:
        print(f"No tokens available for {title}")

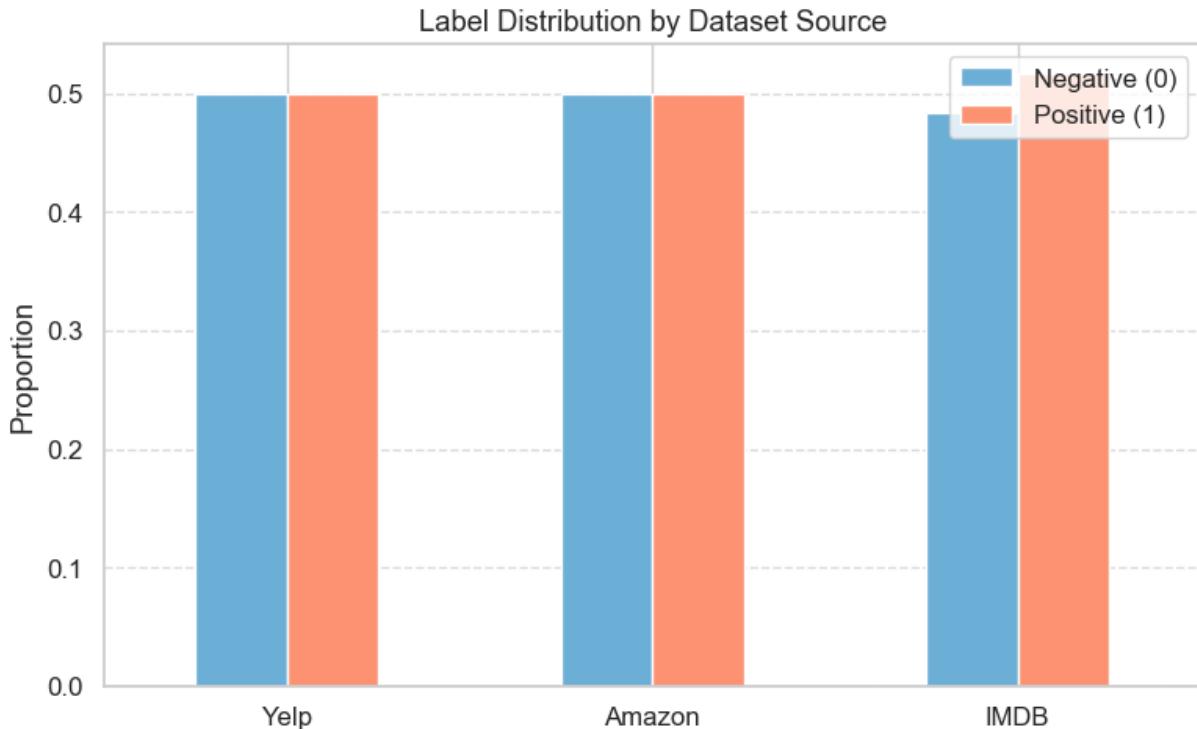
# Plot top words for all labels combined
plot_top_words(df['tokens'], 'Top Words (All Labels)', top_k=25, color="#74c476")

# Plot top words per label
label_colors = ['#6baed6', '#fc9272']
for i, lbl in enumerate(sorted(df['label'].unique())):
    plot_top_words(df.loc[df['label']==lbl, 'tokens'], f'Top Words (Label={lbl})',
                  color=label_colors[i])

# --- Word clouds ---
for i, lbl in enumerate(sorted(df['label'].unique())):
    # Combine all tokens for label into single string
    text_blob = " ".join([" ".join(toks) for toks in df.loc[df['label']==lbl, 'tokens']])
```

```
if text_blob.strip():
    wc = WordCloud(
        width=900, height=450,
        background_color='white',
        colormap='Paired', # color palette
        max_words=150,
        contour_color='black',
        contour_width=1
    ).generate(text_blob)
    plt.figure(figsize=(12,6))
    plt.imshow(wc, interpolation='bilinear') # render word cloud
    plt.axis('off')
    plt.title(f'Word Cloud (Label={lbl})', fontsize=16)
    plt.tight_layout()
    plt.show()
else:
    print(f"No tokens for label {lbl} to build word cloud")

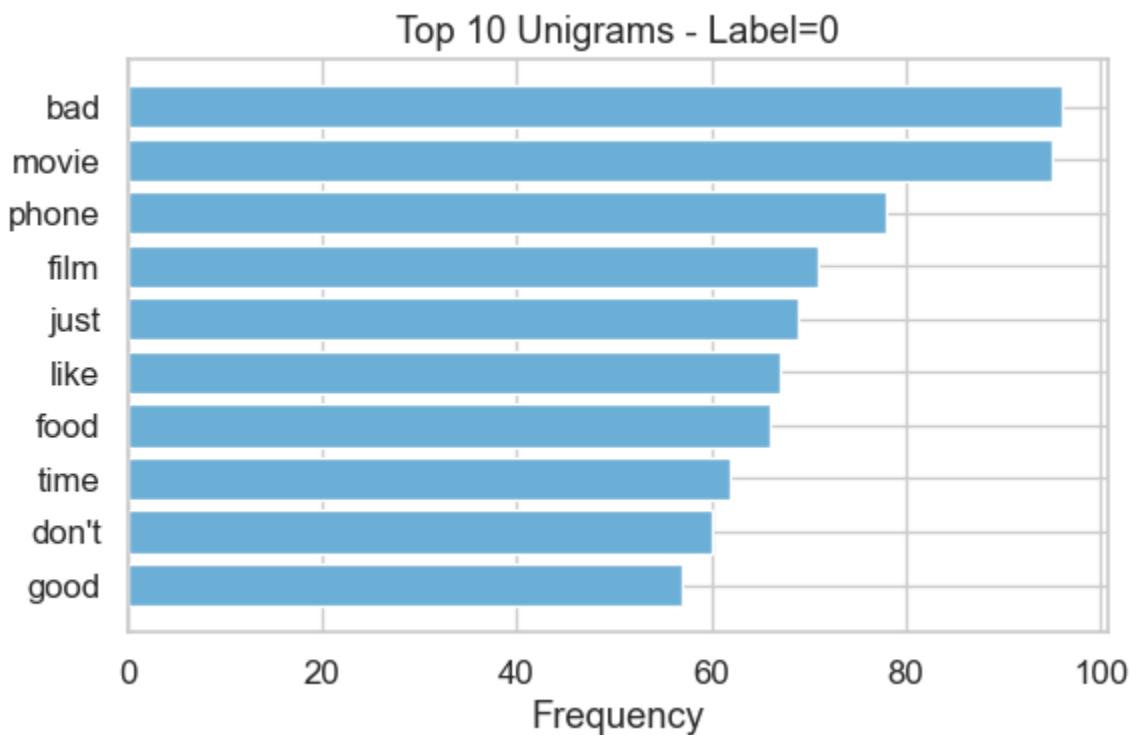
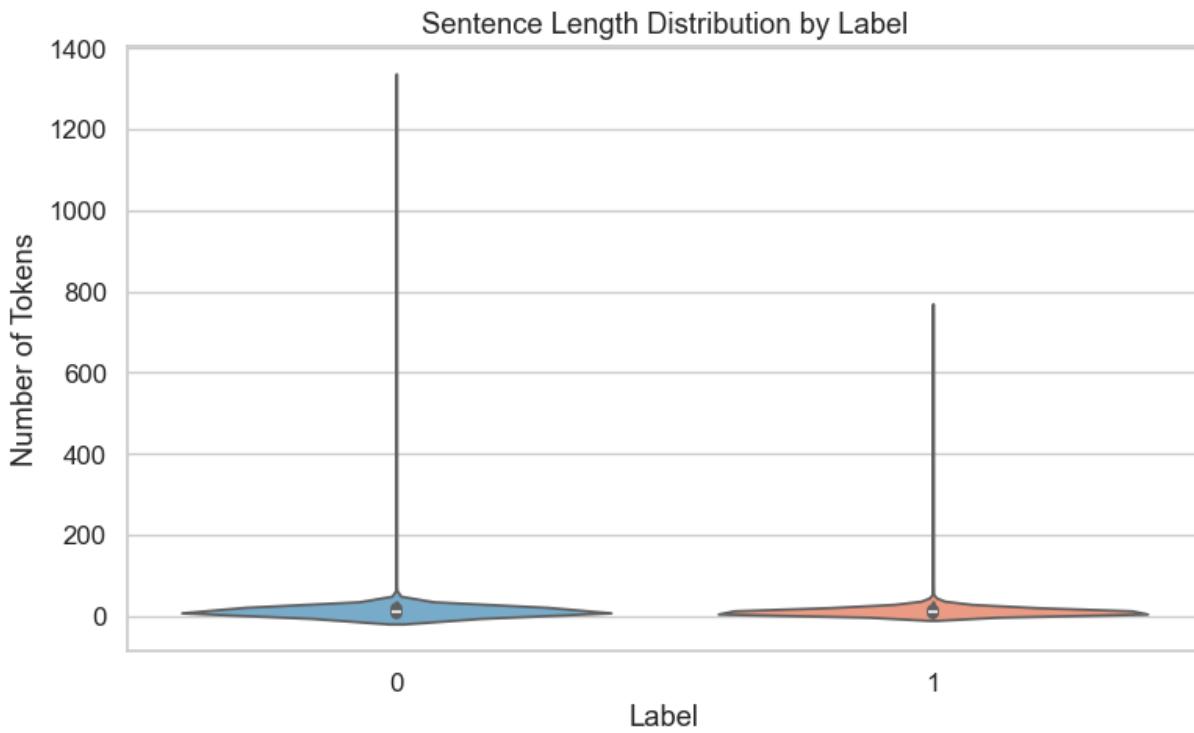
# --- Sentence Length histogram ---
plt.figure(figsize=(8,4))
sns.histplot(df['len_tokens'], bins=30, kde=True, color='#fc9272') # histogram + K
plt.title('Sentence Length Distribution (Cleaned Tokens)', fontsize=14)
plt.xlabel('Number of Tokens', fontsize=12)
plt.ylabel('Frequency', fontsize=12)
plt.tight_layout()
plt.show()
```



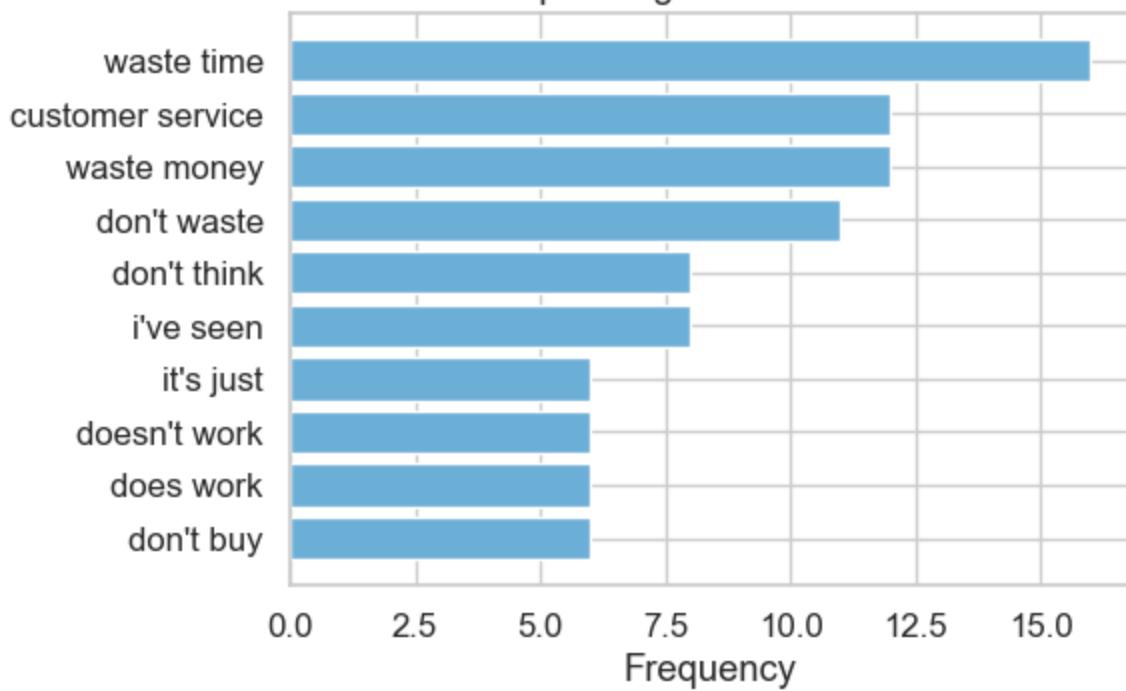
C:\Users\kayam\AppData\Local\Temp\ipykernel\_10724\2476994674.py:29: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14 .0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

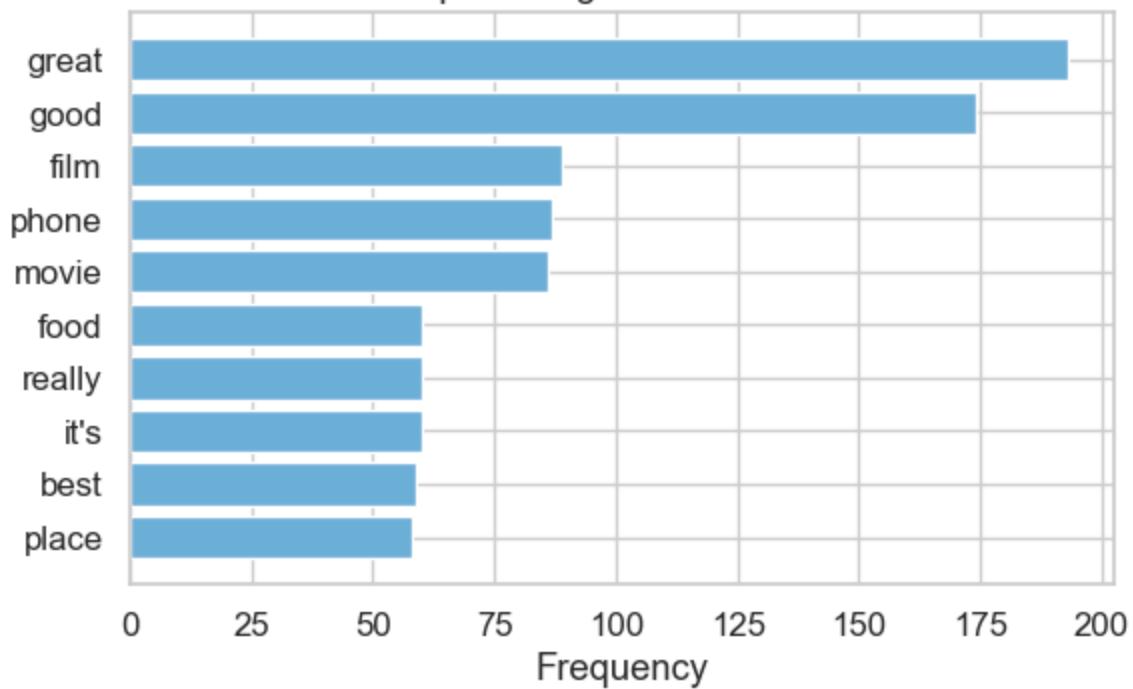
```
sns.violinplot(x='label', y='len_tokens', data=df, palette=['#6baed6', '#fc9272'])
```

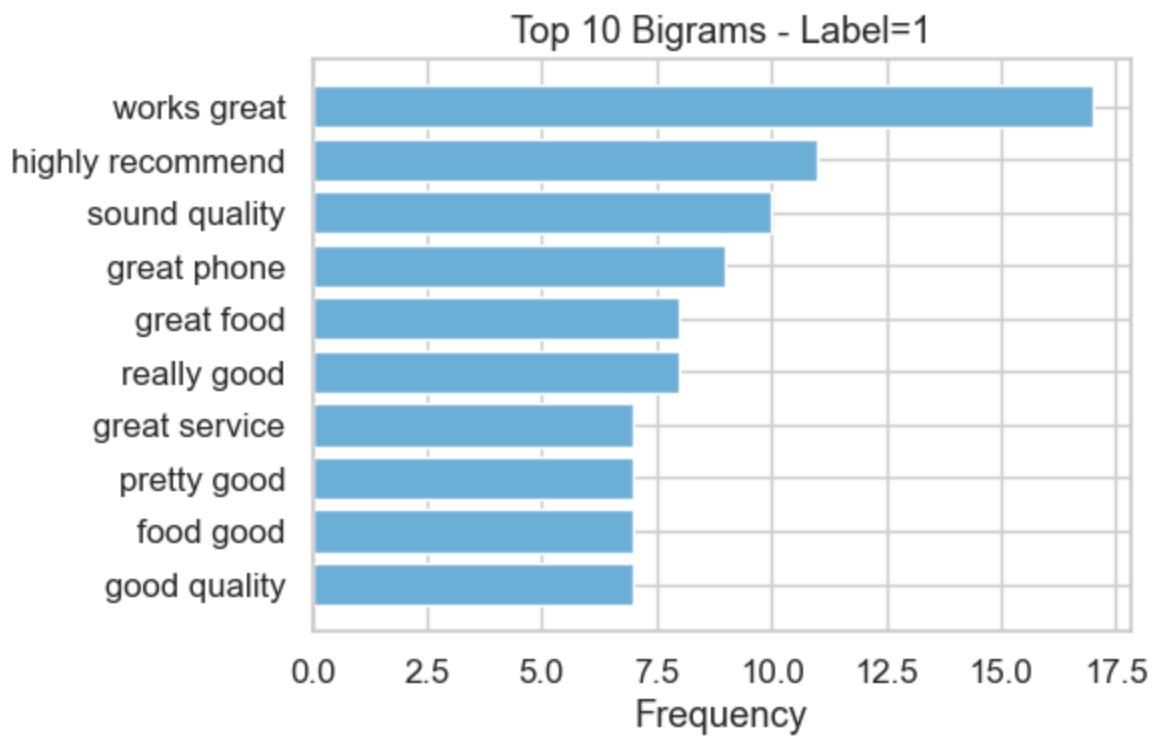


Top 10 Bigrams - Label=0

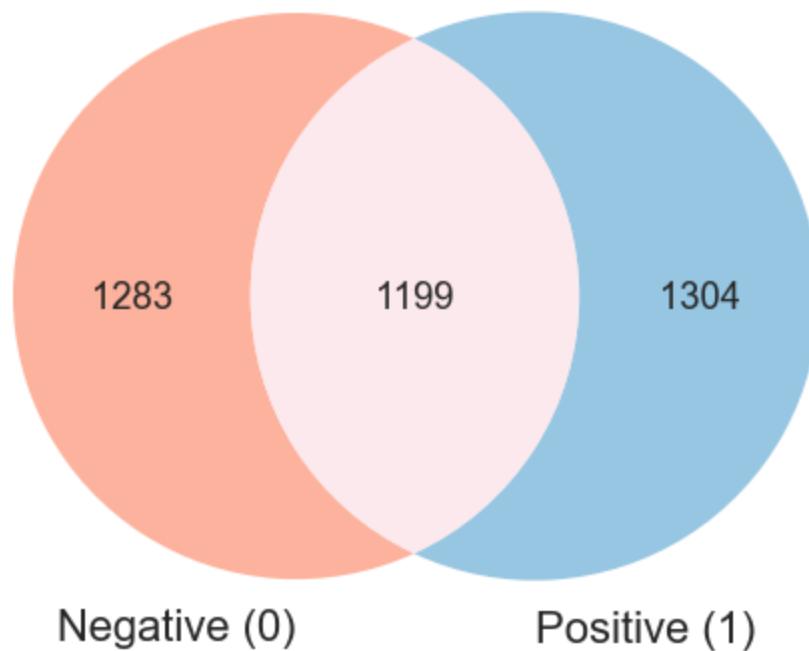


Top 10 Unigrams - Label=1





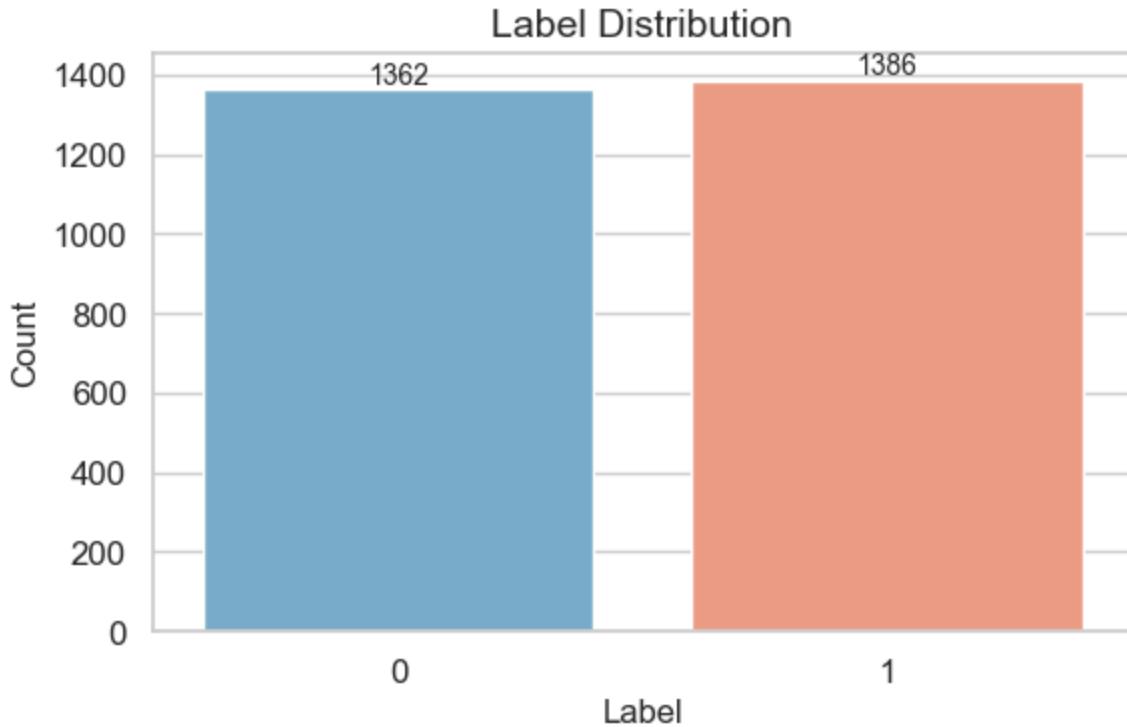
Vocabulary Overlap Between Labels



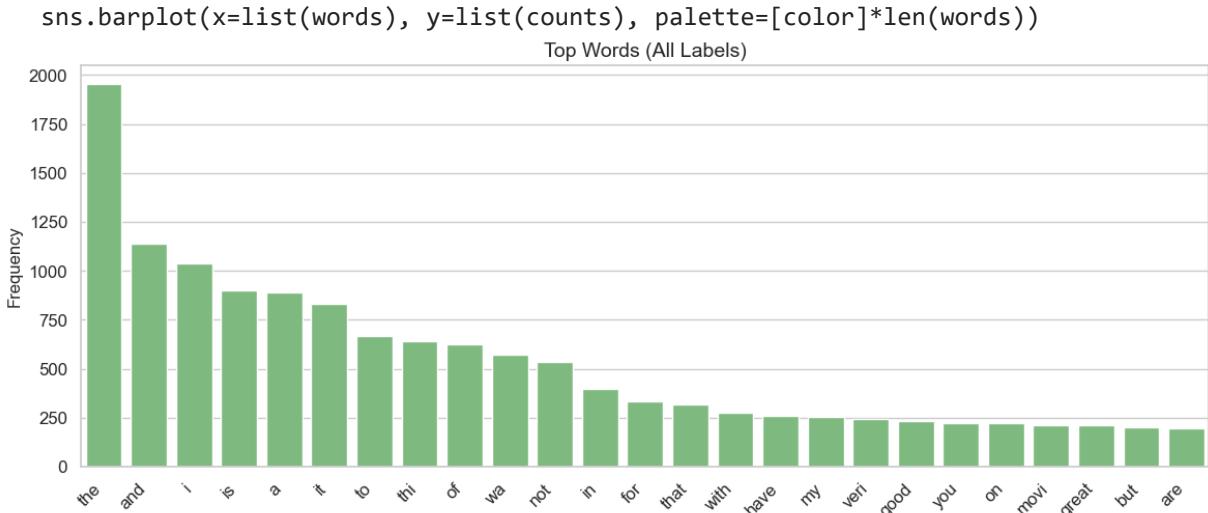
C:\Users\kayam\AppData\Local\Temp\ipykernel\_10724\2476994674.py:83: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14 .0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=label_counts.index, y=label_counts.values, palette=['#6baed6', '#fc9272'])
```

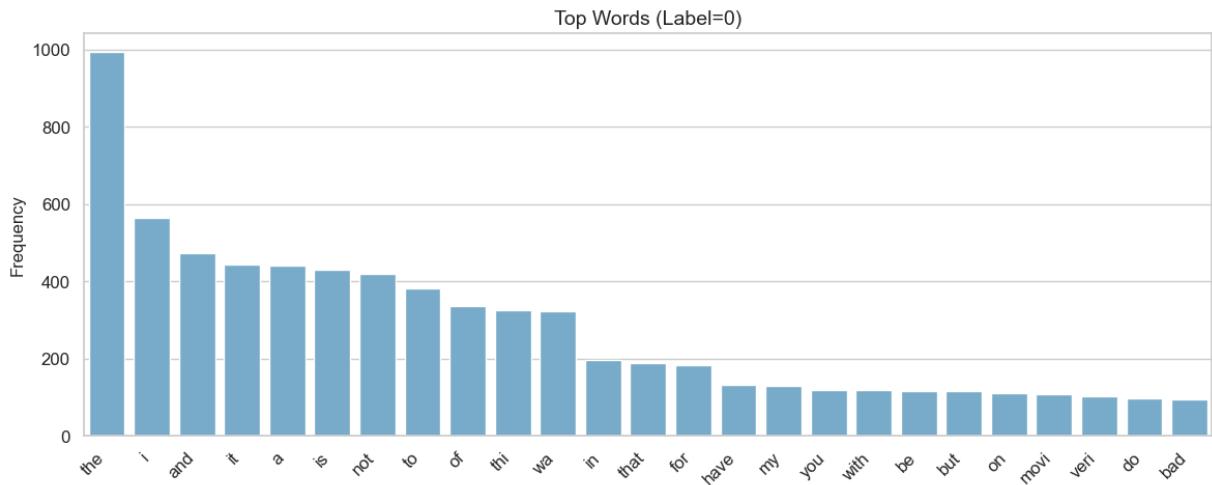


```
C:\Users\kayam\AppData\Local\Temp\ipykernel_10724\2476994674.py:101: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14  
.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```



```
C:\Users\kayam\AppData\Local\Temp\ipykernel_10724\2476994674.py:101: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14  
.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```

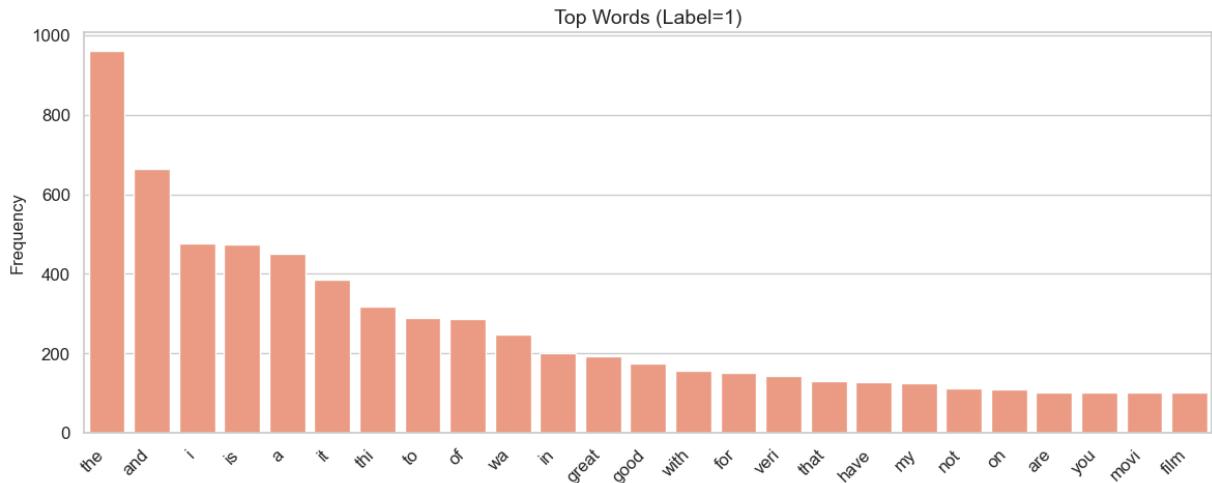
```
sns.barplot(x=list(words), y=list(counts), palette=[color]*len(words))
```



C:\Users\kayam\AppData\Local\Temp\ipykernel\_10724\2476994674.py:101: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=list(words), y=list(counts), palette=[color]*len(words))
```





## 2.3) Train/Validation/Test Split (70/15/15) with CSV Save

```
In [5]: from sklearn.model_selection import train_test_split # for splitting data into tra

# Features (X) include original sentence, cleaned text, and token list
X = df[['sentence','clean_text','tokens']].copy()
# Labels (y) are the sentiment values (0 or 1), ensure integer type
y = df['label'].astype(int)

# --- First split: 70% train, 30% temporary (to be split into val/test) ---
X_train, X_temp, y_train, y_temp = train_test_split(
    X, y, test_size=0.3, stratify=y, random_state=42 # stratify ensures label dist
)

# --- Second split: split temp into 50% validation, 50% test (15% each overall) ---
```

```
X_val, X_test, y_val, y_test = train_test_split(  
    X_temp, y_temp, test_size=0.5, stratify=y_temp, random_state=42  
)  
  
# --- Print shapes to confirm sizes ---  
print("Shapes:")  
print("Train:", X_train.shape, y_train.shape)  
print("Validation:", X_val.shape, y_val.shape)  
print("Test:", X_test.shape, y_test.shape)  
  
import os  
  
# Ensure that a folder exists to save CSVs  
os.makedirs('./data/', exist_ok=True)  
  
# --- Save train/val/test splits as CSV for reproducibility ---  
X_train.join(y_train).to_csv('./data/train_split.csv', index=False) # combine feat  
X_val.join(y_val).to_csv('./data/val_split.csv', index=False)  
X_test.join(y_test).to_csv('./data/test_split.csv', index=False)  
print("Train/Val/Test splits saved to ./data/")
```

Shapes:  
Train: (1923, 3) (1923,)  
Validation: (412, 3) (412,)  
Test: (413, 3) (413,)  
Train/Val/Test splits saved to ./data/

### 3.1 TF-IDF features for classical ML

In [6]:

```
from sklearn.feature_extraction.text import TfidfVectorizer # to convert text to n  
  
# --- Custom stopwords list: remove default stopwords but keep negations (important  
negations = {"not", "no", "never"}  
custom_stopwords = list(ENGLISH_STOP_WORDS.difference(negations))  
  
# Hyperparameters for TF-IDF  
max_features = 5000 # maximum number of features to keep  
ngram_range = (1, 2) # include unigrams and bigrams  
  
# Initialize TF-IDF vectorizer with custom settings  
tfidf_vectorizer = TfidfVectorizer(  
    max_features=max_features,  
    ngram_range=ngram_range,  
    stop_words=custom_stopwords  
)  
  
# Convert train/val/test cleaned text to string type  
X_train_text = X_train['clean_text'].astype(str)  
X_val_text = X_val['clean_text'].astype(str)  
X_test_text = X_test['clean_text'].astype(str)  
  
# Fit TF-IDF on training data, then transform val/test data  
X_train_tfidf = tfidf_vectorizer.fit_transform(X_train_text)  
X_val_tfidf = tfidf_vectorizer.transform(X_val_text)  
X_test_tfidf = tfidf_vectorizer.transform(X_test_text)
```

```
# Print shapes to confirm number of features and samples
print("TF-IDF shapes:")
print("Train:", X_train_tfidf.shape)
print("Val: ", X_val_tfidf.shape)
print("Test: ", X_test_tfidf.shape)

import joblib
# Ensure folder exists to save trained models
os.makedirs('./models/', exist_ok=True)

# Save the fitted TF-IDF vectorizer for later use
joblib.dump(tfidf_vectorizer, './models/tfidf_vectorizer.pkl')
print("TF-IDF vectorizer saved to ./models/tfidf_vectorizer.pkl")
```

```
TF-IDF shapes:
Train: (1923, 5000)
Val:  (412, 5000)
Test: (413, 5000)
TF-IDF vectorizer saved to ./models/tfidf_vectorizer.pkl
```

## 3.2 Tokenization and Padding for Deep Learning Model

```
In [7]: from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

# --- Hyperparameters for deep Learning ---
max_words = 10000 # maximum number of words to keep in tokenizer (top most frequent)
max_len = 100      # maximum sequence length for padding/truncating

# Convert token lists back to strings because Keras tokenizer expects raw text
X_train_text_dl = X_train['tokens'].apply(lambda toks: " ".join(toks))
X_val_text_dl   = X_val['tokens'].apply(lambda toks: " ".join(toks))
X_test_text_dl = X_test['tokens'].apply(lambda toks: " ".join(toks))

# Initialize tokenizer, with special token for out-of-vocabulary words
tokenizer_dl = Tokenizer(num_words=max_words, oov_token=<OOV>)
tokenizer_dl.fit_on_texts(X_train_text_dl) # fit tokenizer only on training data

# Convert text to sequences of integers
X_train_seq_dl = tokenizer_dl.texts_to_sequences(X_train_text_dl)
X_val_seq_dl   = tokenizer_dl.texts_to_sequences(X_val_text_dl)
X_test_seq_dl = tokenizer_dl.texts_to_sequences(X_test_text_dl)

# Pad sequences to the same length for model input
# 'post' padding/truncating means extra zeros or truncation occur at the end
X_train_pad_dl = pad_sequences(X_train_seq_dl, maxlen=max_len, padding='post', truncating='post')
X_val_pad_dl   = pad_sequences(X_val_seq_dl, maxlen=max_len, padding='post', truncating='post')
X_test_pad_dl = pad_sequences(X_test_seq_dl, maxlen=max_len, padding='post', truncating='post')

# Print shapes to confirm
print("Padded sequence shapes (DL):")
print("Train:", X_train_pad_dl.shape)
print("Val: ", X_val_pad_dl.shape)
```

```

print("Test: ", X_test_pad_dl.shape)

# Save tokenizer information for embedding layers or future preprocessing
word_index_dl = tokenizer_dl.word_index # dictionary mapping word -> index
print("Vocabulary size (DL):", len(word_index_dl))

# Save tokenizer object for reuse
joblib.dump(tokenizer_dl, './models/tokenizer_dl.pkl')
print("Tokenizer saved to ./models/tokenizer_dl.pkl")

```

Padded sequence shapes (DL):  
 Train: (1923, 100)  
 Val: (412, 100)  
 Test: (413, 100)  
 Vocabulary size (DL): 3390  
 Tokenizer saved to ./models/tokenizer\_dl.pkl

## 4.1 Classical Model: Logistic Regression with TF-IDF (Improved, Namespaced)

```

In [8]: from sklearn.model_selection import GridSearchCV
        from sklearn.naive_bayes import MultinomialNB
        from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import (classification_report, confusion_matrix, accuracy_score,
                                       precision_recall_curve, auc)
        import seaborn as sns

        # --- Use pre-defined train/val/test labels for classical ML ---
        y_train_cls = y_train
        y_val_cls = y_val
        y_test_cls = y_test

        # --- Naive Bayes model: hyperparameter grid search ---
        nb_params = {'alpha': [0.1, 0.5, 1.0, 2.0]} # smoothing parameter
        nb_grid_cls = GridSearchCV(MultinomialNB(), param_grid=nb_params, cv=5,
                                   scoring='accuracy', n_jobs=-1) # 5-fold CV
        nb_grid_cls.fit(X_train_tfidf, y_train_cls) # fit grid search on training TF-IDF f

        # Print best parameters and CV accuracy for Naive Bayes
        print("Best Naive Bayes params:", nb_grid_cls.best_params_)
        print("Best CV accuracy (NB):", nb_grid_cls.best_score_)

        # --- Logistic Regression model: hyperparameter grid search ---
        log_params = {'C': [0.01, 0.1, 1, 10], 'solver': ['lbfgs', 'liblinear']} # C: regularization parameter
        log_grid_cls = GridSearchCV(LogisticRegression(max_iter=1000), param_grid=log_params,
                                    cv=5, scoring='accuracy', n_jobs=-1)
        log_grid_cls.fit(X_train_tfidf, y_train_cls)

        # Print best parameters and CV accuracy for Logistic Regression
        print("Best Logistic Regression params:", log_grid_cls.best_params_)
        print("Best CV accuracy (LogReg):", log_grid_cls.best_score_)

        # --- Select the better model based on CV accuracy ---
        if log_grid_cls.best_score_ >= nb_grid_cls.best_score_:
            best_model_cls = log_grid_cls.best_estimator_

```

```
    print("Selected model: Logistic Regression")
else:
    best_model_cls = nb_grid_cls.best_estimator_
    print("Selected model: Naive Bayes")

Best Naive Bayes params: {'alpha': 2.0}
Best CV accuracy (NB): 0.8211038961038961
Best Logistic Regression params: {'C': 1, 'solver': 'lbfgs'}
Best CV accuracy (LogReg): 0.8190219155844156
Selected model: Naive Bayes
```

## 4.2 Deep Learning Model: BiLSTM (Improved, Light Hyperparameter Search)

```
In [9]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Bidirectional, Dense, Dropout,
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.regularizers import l2

# -----
# Reproducibility: ensure consistent results
# -----
np.random.seed(42)          # for numpy operations
tf.random.set_seed(42)       # for TensorFlow operations

# -----
# Hyperparameter options (light search to save time)
# -----
embedding_dims_dl = [50, 100]      # size of word embeddings (2 options)
lstm_units_list_dl = [32]          # number of LSTM units (1 option to reduce Load
dropout_rates_dl = [0.3]           # dropout rate after dense Layer
spatial_dropouts_dl = [0.2]         # spatial dropout after embedding Layer
batch_size_dl = 64                 # number of samples per batch
epochs_dl = 12                     # max epochs for training (reduce for speed)

# -----
# Labels as numpy arrays (required by Keras)
# -----
y_train_dl = np.array(y_train)
y_val_dl = np.array(y_val)
y_test_dl = np.array(y_test)

# -----
# Vocabulary size: Limited by max_words and tokenizer
# -----
vocab_size_dl = min(max_words, len(word_index_dl) + 1)

# Variables to track best model during hyperparameter search
best_val_acc_dl = 0.0
best_model_dl = None
best_params_dl = {}

# -----
# Light hyperparameter search: nested loops
```

```
# -----
# for embedding_dim_dl in embedding_dims_dl:
#     for lstm_units_dl in lstm_units_list_dl:
#         for dropout_rate_dl in dropout_rates_dl:
#             for spatial_dropout_dl in spatial_dropouts_dl:

#                 # Print current hyperparameters
#                 print(
#                     f"\nTraining BiLSTM (DL) with emb={embedding_dim_dl}, "
#                     f"lstm={lstm_units_dl}, dropout={dropout_rate_dl}, spatial={spa
#                 )

# -----
# Build BiLSTM model
# -----
model_dl = Sequential([
    # Embedding Layer converts integer sequences to dense vectors
    Embedding(
        input_dim=vocab_size_dl,
        output_dim=embedding_dim_dl,
        input_length=max_len
    ),
    # SpatialDropout1D drops entire embedding dimensions to reduce overfitting
    SpatialDropout1D(spatial_dropout_dl),
    # Bidirectional LSTM layer
    Bidirectional(LSTM(
        lstm_units_dl,
        return_sequences=False,      # Last output only
        dropout=0.3,                # dropout on input connections
        recurrent_dropout=0.3,       # dropout on recurrent connections
        kernel_regularizer=l2(1e-4),   # L2 regularization on weights
        recurrent_regularizer=l2(1e-4) # L2 regularization on recurrent weights
    )),
    # Dense hidden layer
    Dense(32, activation='relu', kernel_regularizer=l2(1e-4)),
    Dropout(dropout_rate_dl),      # dropout for regularization
    # Output layer for binary classification
    Dense(1, activation='sigmoid')
])

# -----
# Compile model
# -----
model_dl.compile(
    loss='binary_crossentropy',    # suitable for binary classification
    optimizer='adam',              # adaptive optimizer
    metrics=['accuracy']          # track accuracy during training
)

# -----
# Early stopping: stop if no improvement in validation loss
# -----
early_stop_dl = EarlyStopping(
    monitor='val_loss',           # monitor validation loss
    patience=2,                  # wait for 2 epochs before stopping
    restore_best_weights=True      # keep best weights
```

```
)\n\n# -----\n# Train model\n#\nhistory_dl = model_dl.fit(\n    X_train_pad_dl, y_train_dl,\n    validation_data=(X_val_pad_dl, y_val_dl),\n    epochs=epochs_dl,\n    batch_size=batch_size_dl,\n    callbacks=[early_stop_dl], # apply early stopping\n    verbose=0                 # silent output\n)\n\n# -----\n# Track maximum validation accuracy\n#\nval_acc_dl = max(history_dl.history['val_accuracy'])\nprint(f"Max val accuracy (DL): {val_acc_dl:.3f}")\n\n# -----\n# Keep best model and hyperparameters\n#\nif val_acc_dl > best_val_acc_dl:\n    best_val_acc_dl = val_acc_dl\n    best_model_dl = model_dl\n    best_params_dl = {\n        'embedding_dim': embedding_dim_dl,\n        'lstm_units': lstm_units_dl,\n        'dropout_rate': dropout_rate_dl,\n        'spatial_dropout': spatial_dropout_dl\n    }\n\n# -----\n# Print best hyperparameters and validation accuracy\n#\nprint("\nBest hyperparameters found for DL model:")\nprint(best_params_dl)\nprint(f"Best Validation Accuracy (DL): {best_val_acc_dl:.3f}")\n\n# -----\n# Save best model to disk\n#\nos.makedirs('./models/', exist_ok=True) # ensure folder exists\nbest_model_dl.save('./models/bilstm_model_best_dl.keras')\nprint("Best BiLSTM DL model saved to ./models/bilstm_model_best_dl.keras")
```

Training BiLSTM (DL) with emb=50, lstm=32, dropout=0.3, spatial=0.2

C:\Users\kayam\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\core\embedding.py:97: UserWarning: Argument `input\_length` is deprecated. Just remove it.

```
warnings.warn(
```

```
Max val accuracy (DL): 0.833
```

```
Training BiLSTM (DL) with emb=100, lstm=32, dropout=0.3, spatial=0.2
Max val accuracy (DL): 0.825
```

```
Best hyperparameters found for DL model:
```

```
{'embedding_dim': 50, 'lstm_units': 32, 'dropout_rate': 0.3, 'spatial_dropout': 0.2}
Best Validation Accuracy (DL): 0.833
Best BiLSTM DL model saved to ./models/bilstm_model_best_dl.keras
```

## 5.1 Plots - Classical Machine Learning Model

```
In [10]: # --- Evaluate on validation set ---
# Predict Labels on validation set using the best classical ML model
y_val_pred_cls = best_model_cls.predict(X_val_tfidf)
print("Validation classification report CLASSICAL:")
# Print detailed metrics: precision, recall, f1-score for each class
print(classification_report(y_val_cls, y_val_pred_cls, digits=3))

# --- Evaluate on test set ---
# Predict Labels on the test set
y_test_pred_cls = best_model_cls.predict(X_test_tfidf)
print("Test classification report CLASSICAL:")
# Show metrics on test set
print(classification_report(y_test_cls, y_test_pred_cls, digits=3))

# Compute overall test accuracy
test_acc_cls = accuracy_score(y_test_cls, y_test_pred_cls)
print(f"Test Accuracy (Classical ML): {test_acc_cls:.3f}")

# --- Confusion Matrix ---
# Confusion matrix: shows True vs Predicted counts
cm_cls = confusion_matrix(y_test_cls, y_test_pred_cls)
plt.figure(figsize=(4,3))
sns.heatmap(cm_cls, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Pred 0', 'Pred 1'],
            yticklabels=['True 0', 'True 1'])
plt.title('Confusion Matrix - Best Classical ML Model (Test)')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.tight_layout()
plt.show()

# --- ROC and PR Curves ---
# Get probability estimates for positive class
y_test_proba_cls = best_model_cls.predict_proba(X_test_tfidf)[:,1]

# ROC Curve
fpr_cls, tpr_cls, _ = roc_curve(y_test_cls, y_test_proba_cls)
roc_auc_cls = roc_auc_score(y_test_cls, y_test_proba_cls)

plt.figure(figsize=(6,5))
plt.plot(fpr_cls, tpr_cls, label=f'AUC = {roc_auc_cls:.3f}')
plt.plot([0,1], [0,1], 'k--') # random classifier line
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Best Classical ML Model')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

# Precision-Recall Curve
precision_cls, recall_cls, _ = precision_recall_curve(y_test_cls, y_test_proba_cls)
pr_auc_cls = auc(recall_cls, precision_cls)

plt.figure(figsize=(6,5))
plt.plot(recall_cls, precision_cls, label=f'PR AUC = {pr_auc_cls:.3f}')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve - Best Classical ML Model')
plt.legend(loc='lower left')
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Top TF-IDF Features (Logistic Regression only) ---
# Only applicable if the model has coefficients (like Logistic Regression)
if hasattr(best_model_cls, 'coef_'):
    feature_names = tfidf_vectorizer.get_feature_names_out()
    coefs_cls = best_model_cls.coef_[0]
    # Top 10 positive features (most indicative of positive sentiment)
    top_pos_idx_cls = coefs_cls.argsort()[-10:][::-1]
    # Top 10 negative features (most indicative of negative sentiment)
    top_neg_idx_cls = coefs_cls.argsort()[:10]

    print("\nTop 10 Positive Features (words indicating positive sentiment):")
    for i in top_pos_idx_cls:
        print(f"{feature_names[i]}: {coefs_cls[i]:.3f}")

    print("\nTop 10 Negative Features (words indicating negative sentiment):")
    for i in top_neg_idx_cls:
        print(f"{feature_names[i]}: {coefs_cls[i]:.3f}")
```

## Validation classification report CLASSICAL:

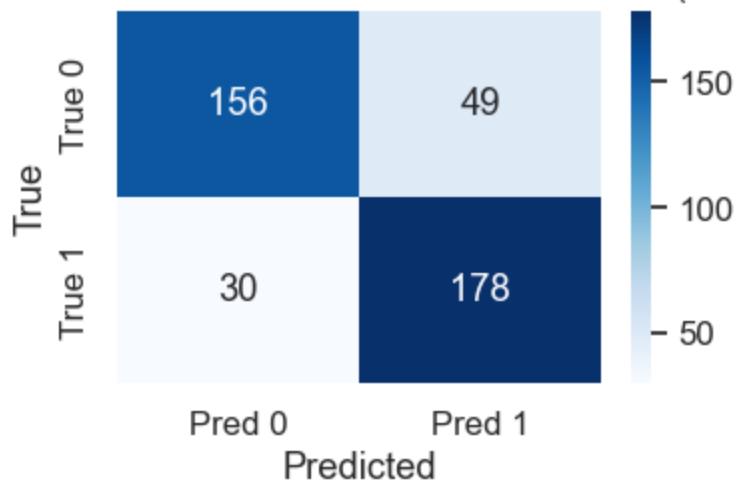
	precision	recall	f1-score	support
0	0.782	0.828	0.805	204
1	0.821	0.774	0.797	208
accuracy			0.801	412
macro avg	0.802	0.801	0.801	412
weighted avg	0.802	0.801	0.801	412

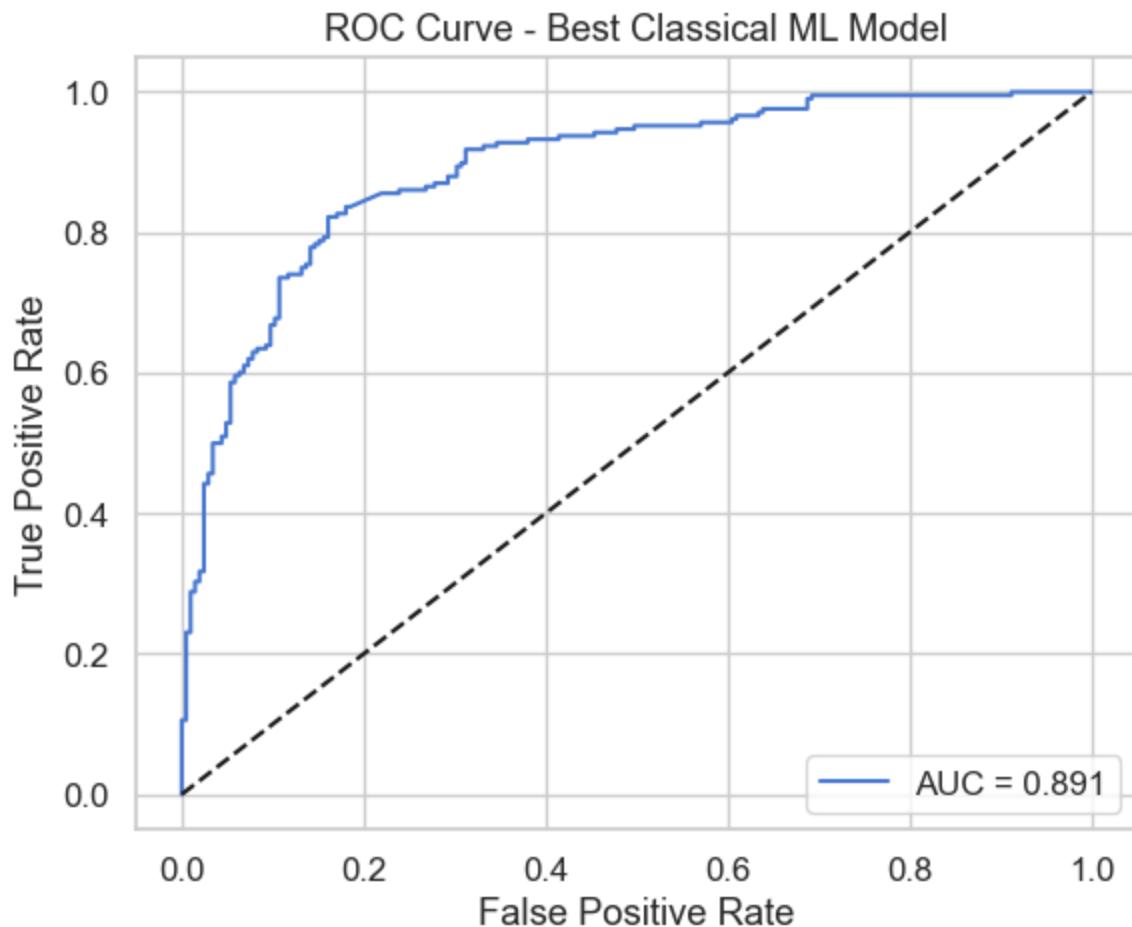
## Test classification report CLASSICAL:

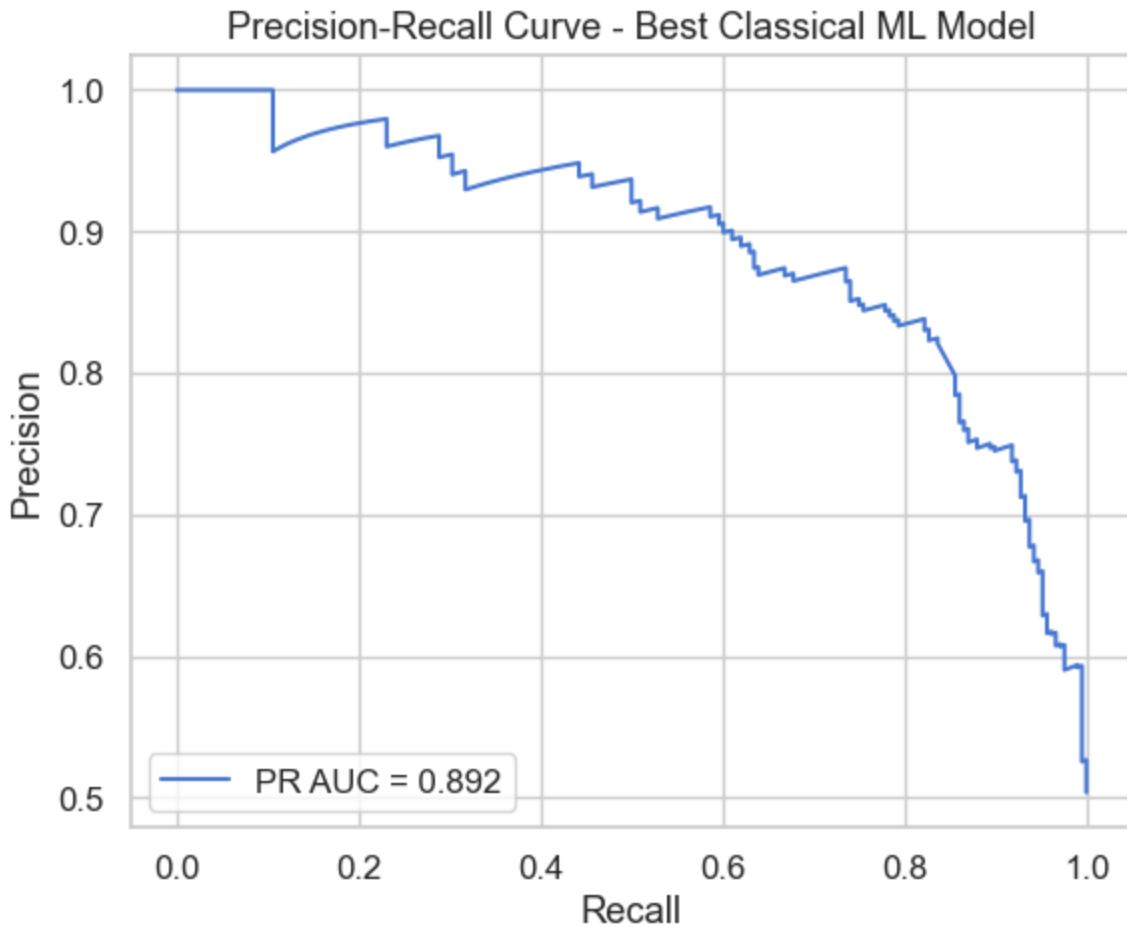
	precision	recall	f1-score	support
0	0.839	0.761	0.798	205
1	0.784	0.856	0.818	208
accuracy			0.809	413
macro avg	0.811	0.808	0.808	413
weighted avg	0.811	0.809	0.808	413

Test Accuracy (Classical ML): 0.809

Confusion Matrix - Best Classical ML Model (Test)





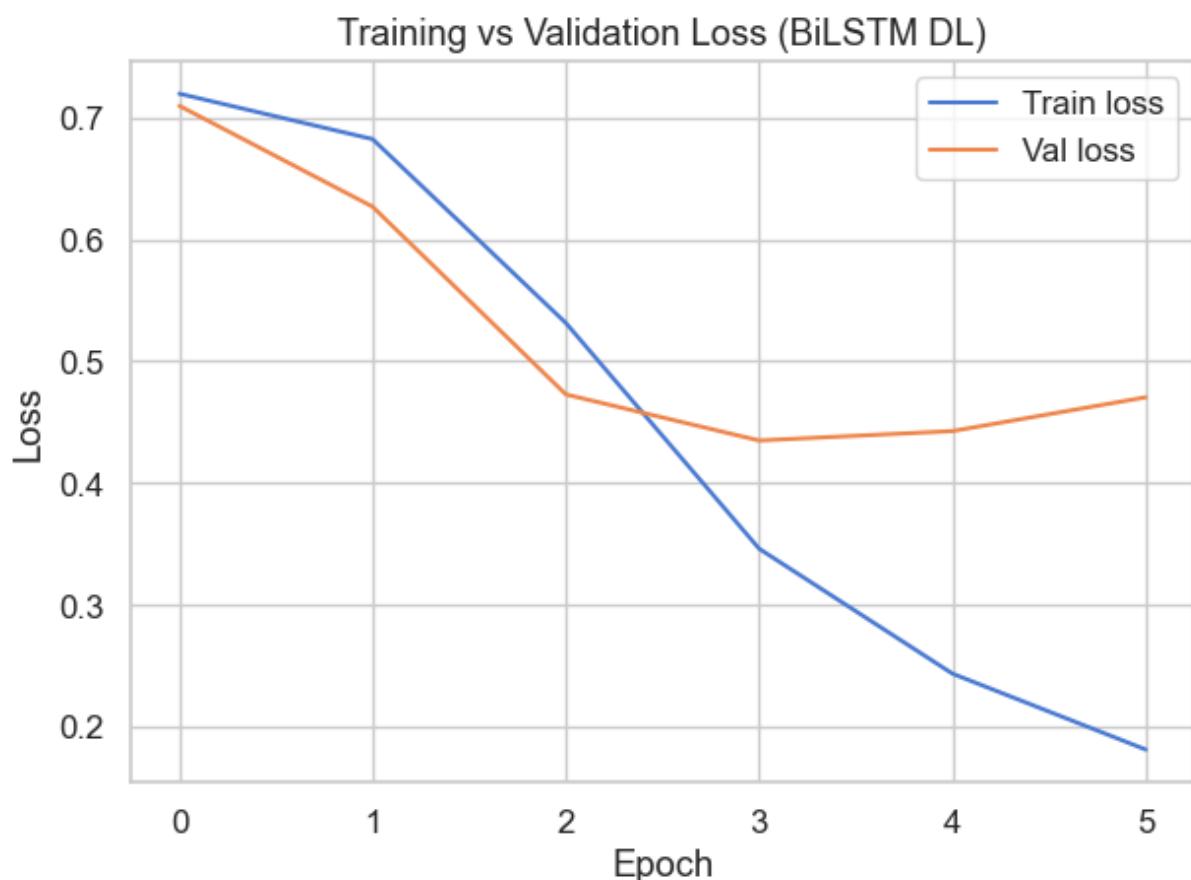
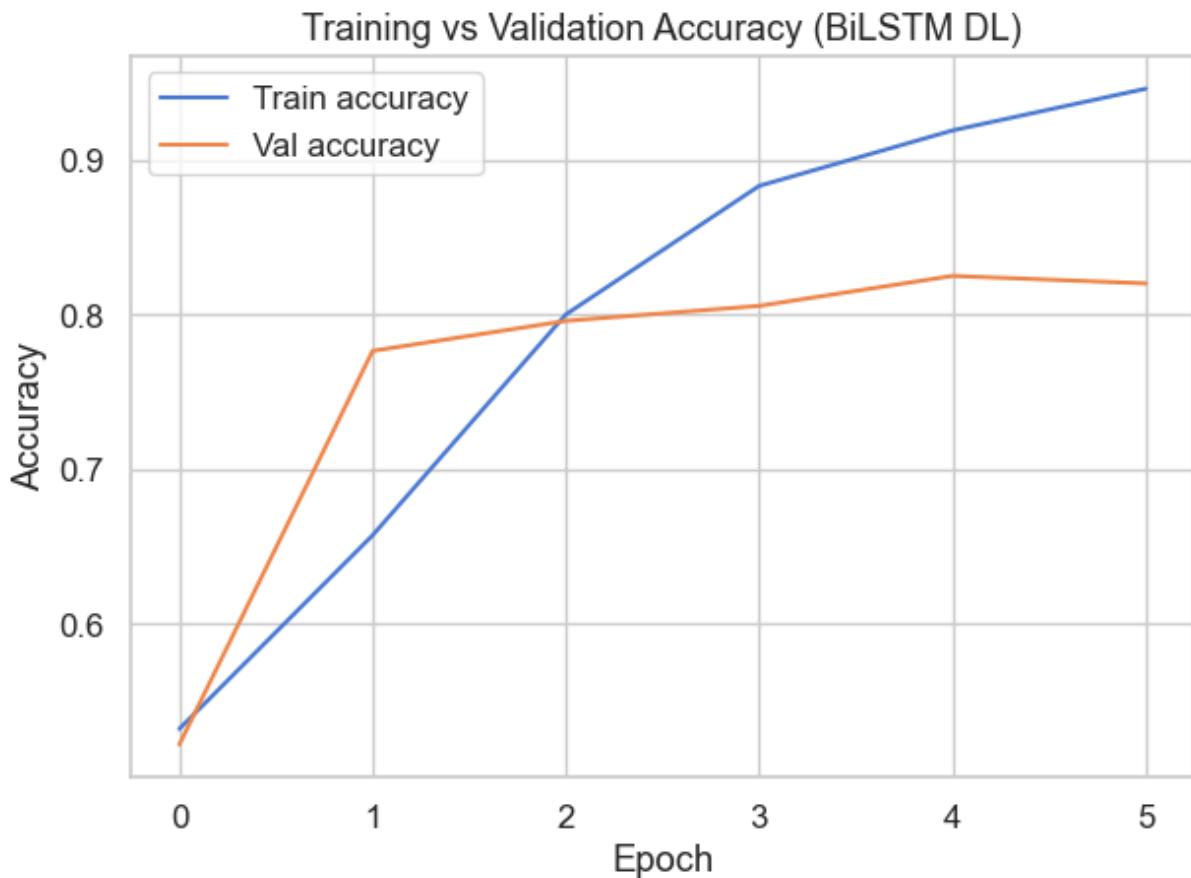


## 5.2 Plot Training Curves (DL)

```
In [11]: # Accuracy Curve
# Plot training and validation accuracy over epochs for BiLSTM model
plt.figure()
plt.plot(history_dl.history['accuracy'], label='Train accuracy')
plt.plot(history_dl.history['val_accuracy'], label='Val accuracy')
plt.title('Training vs Validation Accuracy (BiLSTM DL)')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# Loss Curve
# Plot training and validation loss over epochs for BiLSTM model
plt.figure()
plt.plot(history_dl.history['loss'], label='Train loss')
plt.plot(history_dl.history['val_loss'], label='Val loss')
plt.title('Training vs Validation Loss (BiLSTM DL)')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
```

```
plt.tight_layout()  
plt.show()
```



## 5.3 Evaluation on Test Set (DL)

```
In [12]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

# --- Evaluate overall test Loss and accuracy for BiLSTM DL model ---
test_loss_dl, test_acc_dl = best_model_dl.evaluate(X_test_pad_dl, y_test_dl, verbose=0)
print(f"Test Loss (BiLSTM DL): {test_loss_dl:.3f}")
print(f"Test Accuracy (BiLSTM DL): {test_acc_dl:.3f}")

# --- Predicted probabilities & class labels ---
# Predict probabilities for the positive class (output of sigmoid)
y_test_proba_dl = best_model_dl.predict(X_test_pad_dl).ravel()
# Convert probabilities to binary labels (threshold = 0.5)
y_test_pred_dl = (y_test_proba_dl >= 0.5).astype(int)

# --- Classification report ---
# Show precision, recall, F1-score per class
print("\nClassification report (BiLSTM DL):")
print(classification_report(y_test_dl, y_test_pred_dl, digits=3))

# --- Confusion matrix ---
# Visualize True vs Predicted labels counts
cm_dl = confusion_matrix(y_test_dl, y_test_pred_dl)
plt.figure(figsize=(4,3))
sns.heatmap(cm_dl, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Pred 0', 'Pred 1'],
            yticklabels=['True 0', 'True 1'])
plt.title('Confusion Matrix - BiLSTM DL (Test)')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.tight_layout()
plt.show()

# --- Additional evaluation metrics ---
from sklearn.metrics import roc_curve, roc_auc_score, precision_recall_curve, auc,

# ROC Curve & AUC
fpr_dl, tpr_dl, _ = roc_curve(y_test_dl, y_test_proba_dl)
roc_auc_dl = roc_auc_score(y_test_dl, y_test_proba_dl)

plt.figure(figsize=(6,5))
plt.plot(fpr_dl, tpr_dl, label=f'AUC = {roc_auc_dl:.3f}')
plt.plot([0,1], [0,1], 'k--') # diagonal = random classifier
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - BiLSTM DL')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

# Precision-Recall Curve & AUC
precision_dl_vals, recall_dl_vals, _ = precision_recall_curve(y_test_dl, y_test_proba_dl)
pr_auc_dl = auc(recall_dl_vals, precision_dl_vals)
```

```
plt.figure(figsize=(6,5))
plt.plot(recall_dl_vals, precision_dl_vals, label=f'PR AUC = {pr_auc_dl:.3f}')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve - BiLSTM DL')
plt.legend(loc='lower left')
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Token Importance (Simple Interpretability Proxy) ---
# Extract embedding weights (word vectors) from first layer
embedding_layer_dl = best_model_dl.layers[0]
embedding_weights_dl = embedding_layer_dl.get_weights()[0] # shape: (vocab_size, e

# Compute L2 norm for each word vector as a measure of "importance"
token_importance_dl = np.linalg.norm(embedding_weights_dl, axis=1)

# Reverse Lookup: index → word
index_word_dl = {v: k for k, v in tokenizer_dl.word_index.items()}

# Top 10 most influential tokens based on L2 norm
top_tokens_idx_dl = token_importance_dl.argsort()[-10:][::-1]
top_tokens_dl = [index_word_dl.get(i, "<OOV>") for i in top_tokens_idx_dl]

print("\nTop 10 influential tokens (BiLSTM DL, L2 norm):")
for token, score in zip(top_tokens_dl, token_importance_dl[top_tokens_idx_dl]):
    print(f"{token}: {score:.3f}")

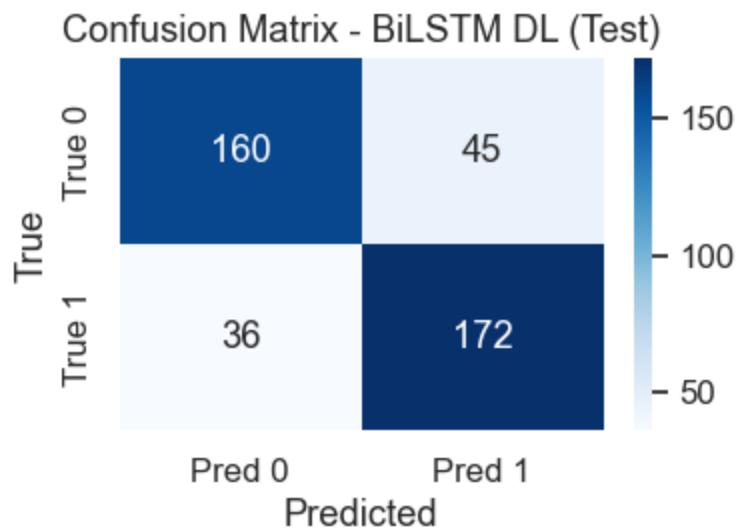
# --- Final Metrics (DL) ---
# Compute standard evaluation metrics
acc_dl = accuracy_score(y_test_dl, y_test_pred_dl)
prec_dl = precision_score(y_test_dl, y_test_pred_dl)
recall_dl = recall_score(y_test_dl, y_test_pred_dl)
f1_dl = f1_score(y_test_dl, y_test_pred_dl)

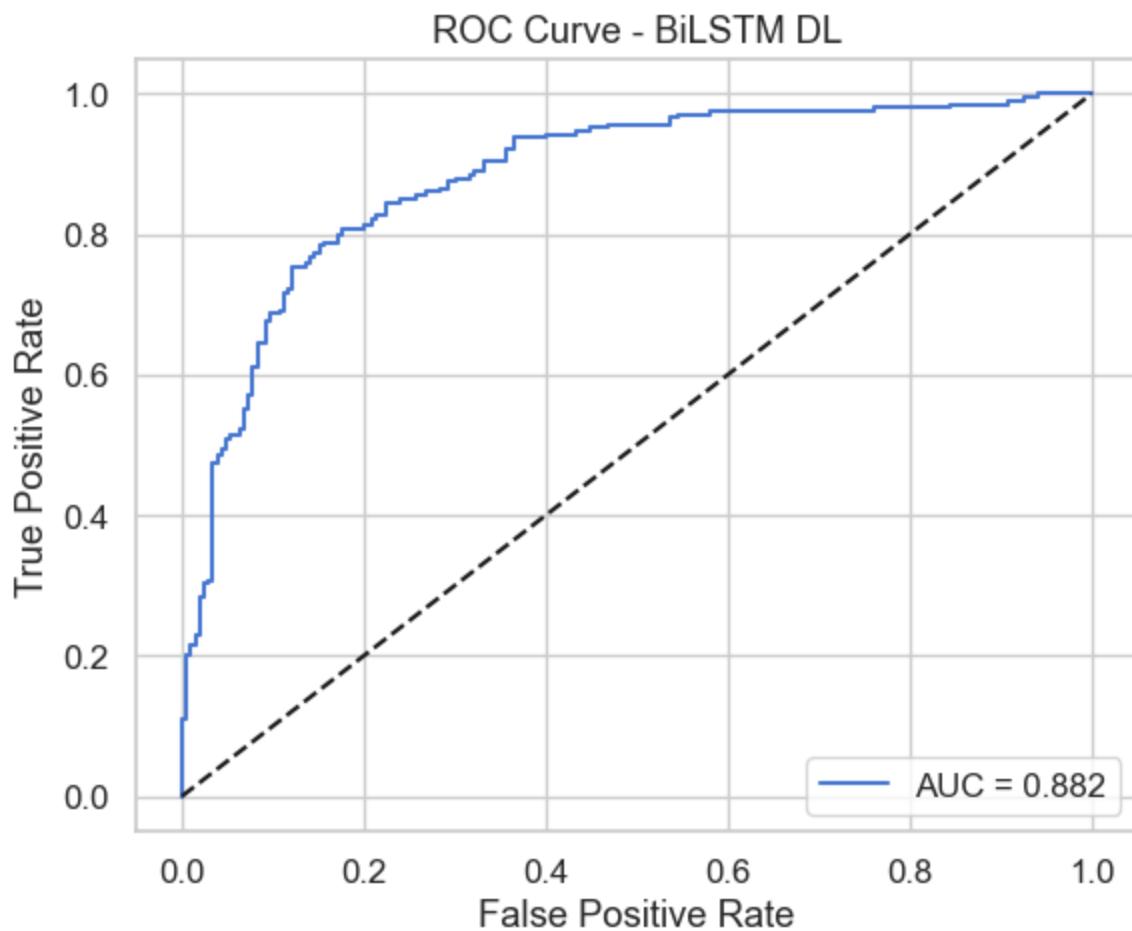
# Print final metrics
print(f"\nBiLSTM DL Test Accuracy: {acc_dl:.3f}")
print(f"BiLSTM DL Precision: {prec_dl:.3f}")
print(f"BiLSTM DL Recall: {recall_dl:.3f}")
print(f"BiLSTM DL F1-score: {f1_dl:.3f}")
```

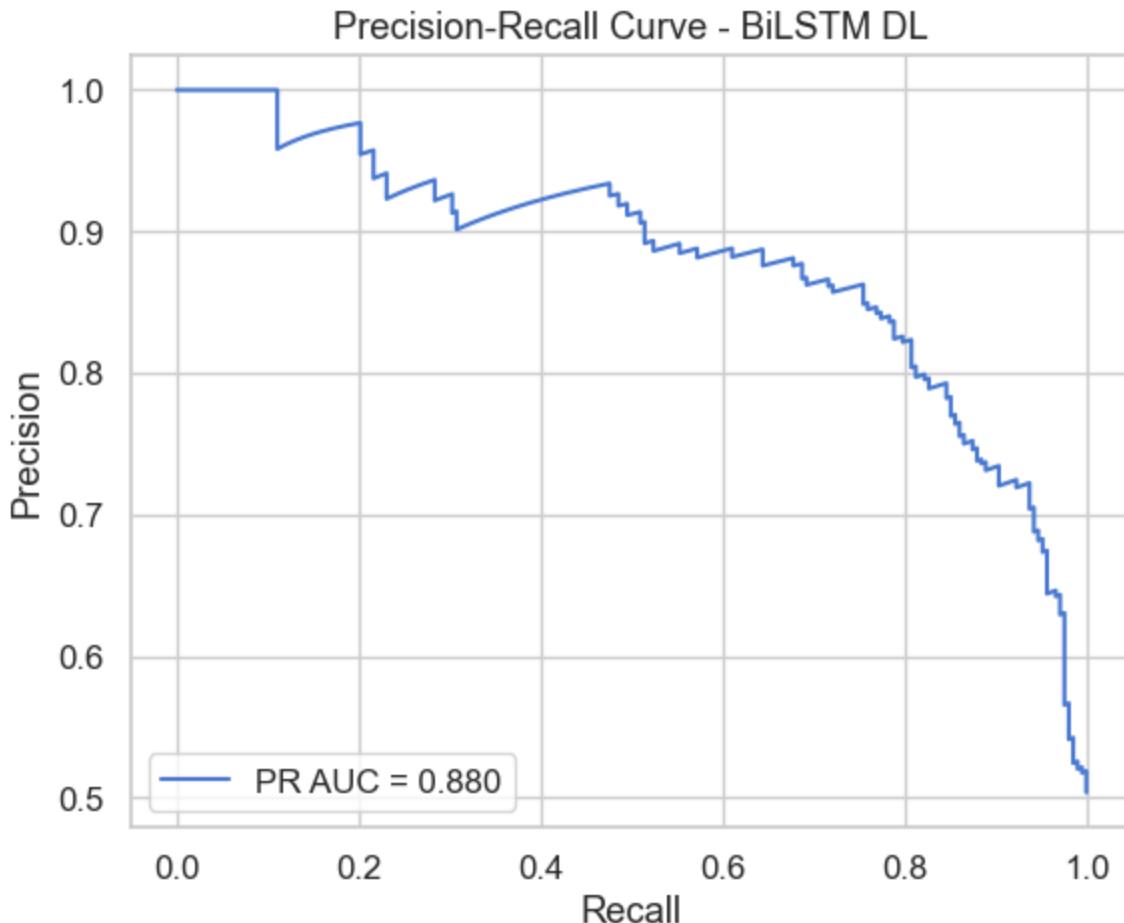
Test Loss (BiLSTM DL): 0.452  
Test Accuracy (BiLSTM DL): 0.804  
**13/13** ————— **4s** 158ms/step

Classification report (BiLSTM DL):

	precision	recall	f1-score	support
0	0.816	0.780	0.798	205
1	0.793	0.827	0.809	208
accuracy			0.804	413
macro avg	0.804	0.804	0.804	413
weighted avg	0.804	0.804	0.804	413







Top 10 influential tokens (BiLSTM DL, L2 norm):

great: 1.054  
not: 0.839  
disappoint: 0.749  
excel: 0.717  
love: 0.699  
bad: 0.688  
nice: 0.686  
wast: 0.677  
worst: 0.640  
good: 0.640

BiLSTM DL Test Accuracy: 0.804  
BiLSTM DL Precision: 0.793  
BiLSTM DL Recall: 0.827  
BiLSTM DL F1-score: 0.809

## 5.4 Combined Model Comparison

```
In [13]: import pandas as pd
import matplotlib.pyplot as plt

# --- Collect Metrics for both models ---
# Create a dictionary to store key evaluation metrics for each model
metrics = {
    "Model": ["Classical ML", "BiLSTM DL"], # Model names
```

```
"Accuracy": [test_acc_cls, acc_dl], # Overall accuracy on test set
"Precision": [precision_score(y_test_cls, y_test_pred_cls),
               prec_dl], # How many predicted positives were correct
"Recall": [recall_score(y_test_cls, y_test_pred_cls),
            recall_dl], # How many true positives were captured
"F1-score": [f1_score(y_test_cls, y_test_pred_cls),
              f1_dl], # Harmonic mean of precision & recall
"ROC-AUC": [roc_auc_cls, roc_auc_dl], # Area under ROC curve
"PR-AUC": [pr_auc_cls, pr_auc_dl] # Area under precision-recall curve
}

# Convert dictionary into a pandas DataFrame for easy tabular display
df_metrics = pd.DataFrame(metrics)
# This DataFrame will be printed and visualized in plots
```

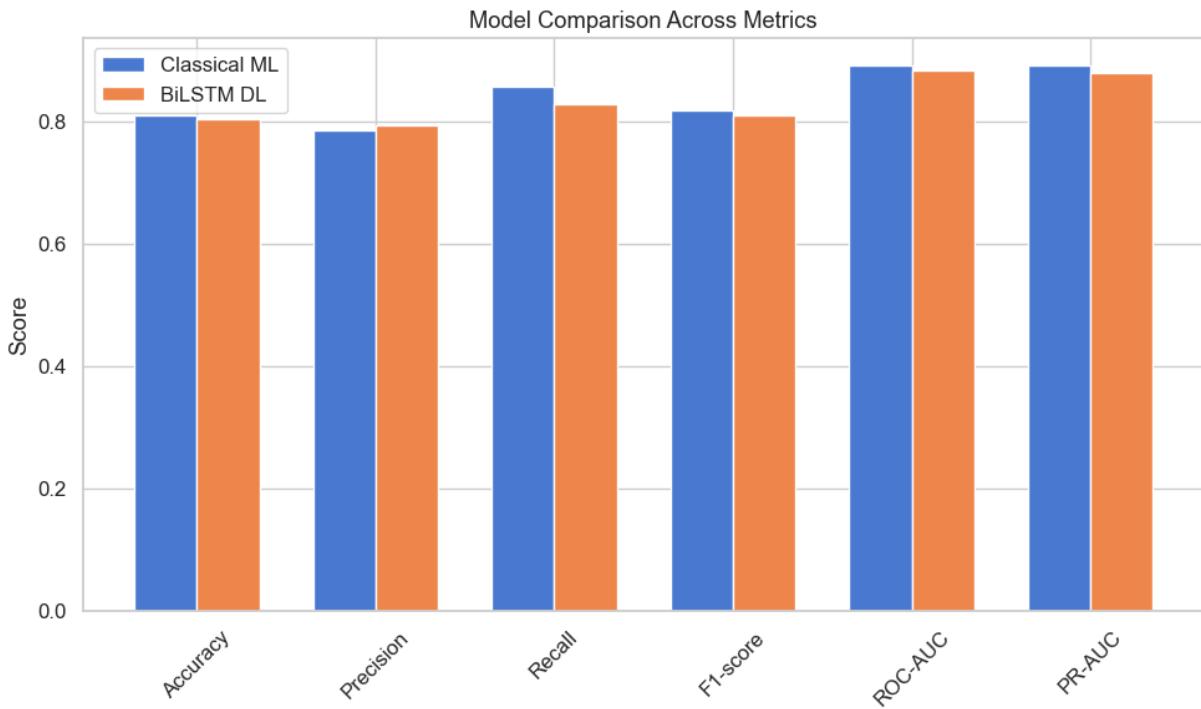
## 5.5 Bar Plot — Metric Comparison

```
In [14]: plt.figure(figsize=(10,6)) # Set figure size

bar_width = 0.35 # Width of each bar
x = np.arange(len(df_metrics.columns[1:])) # Positions on x-axis for metrics (excluding accuracy)

# Plot bars for Classical ML metrics (shifted left)
plt.bar(x - bar_width/2, df_metrics.iloc[0,1:], width=bar_width, label="Classical M")
# Plot bars for BiLSTM DL metrics (shifted right)
plt.bar(x + bar_width/2, df_metrics.iloc[1,1:], width=bar_width, label="BiLSTM DL")

# Set x-axis tick labels as metric names
plt.xticks(x, df_metrics.columns[1:], rotation=45)
plt.ylabel("Score") # Label y-axis
plt.title("Model Comparison Across Metrics") # Add title
plt.legend() # Show legend
plt.tight_layout() # Adjust layout to fit everything
plt.grid(True, axis='y') # Add horizontal grid lines for easier reading
plt.show() # Display the plot
```

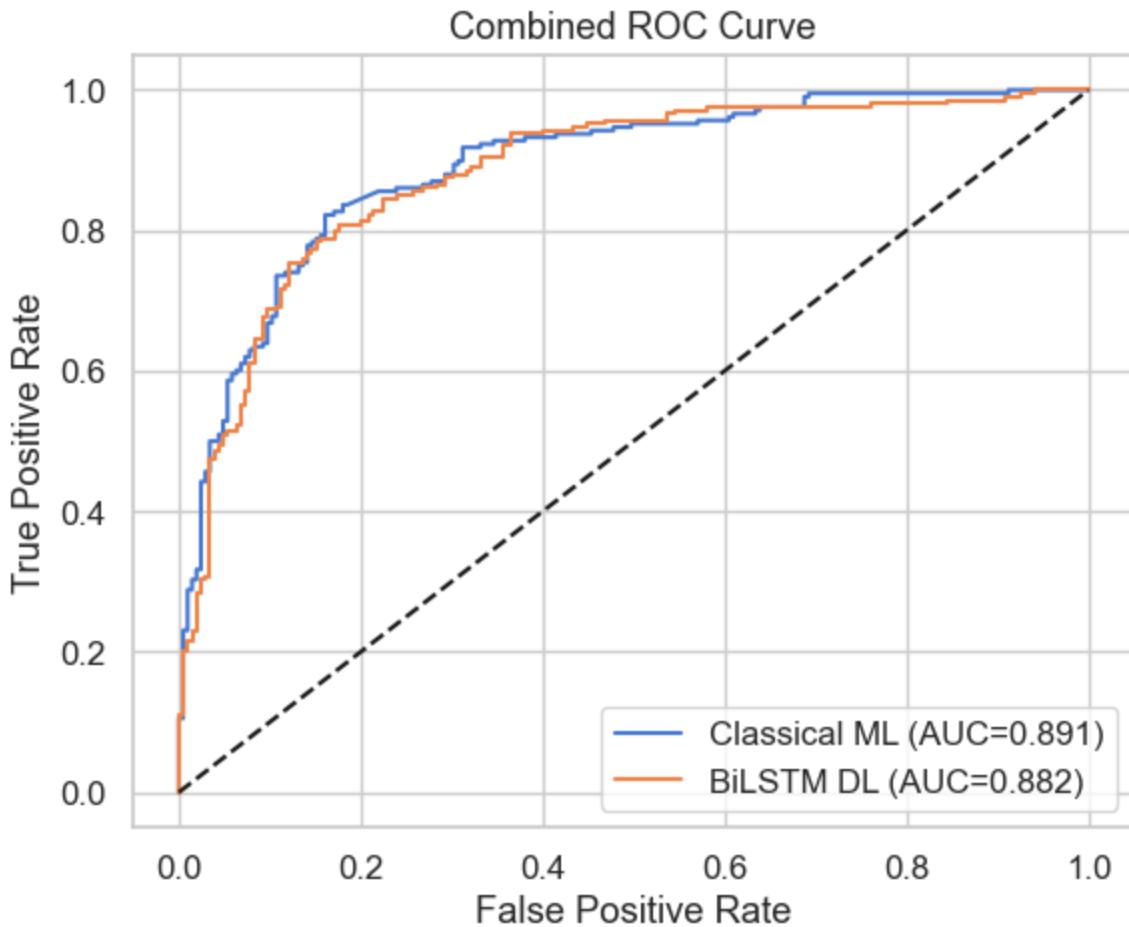


## 5.6 Combined ROC Curve

```
In [15]: plt.figure(figsize=(6,5))

# Plot ROC curve for Classical ML
plt.plot(fpr_cls, tpr_cls, label=f"Classical ML (AUC={roc_auc_cls:.3f})")
# Plot ROC curve for BiLSTM DL
plt.plot(fpr_dl, tpr_dl, label=f"BiLSTM DL (AUC={roc_auc_dl:.3f})")
# Add diagonal line for random classifier reference
plt.plot([0,1], [0,1], 'k--')

# Labels and title
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Combined ROC Curve")
plt.grid(True) # Add grid
plt.legend() # Show legend
plt.tight_layout()
plt.show()
```

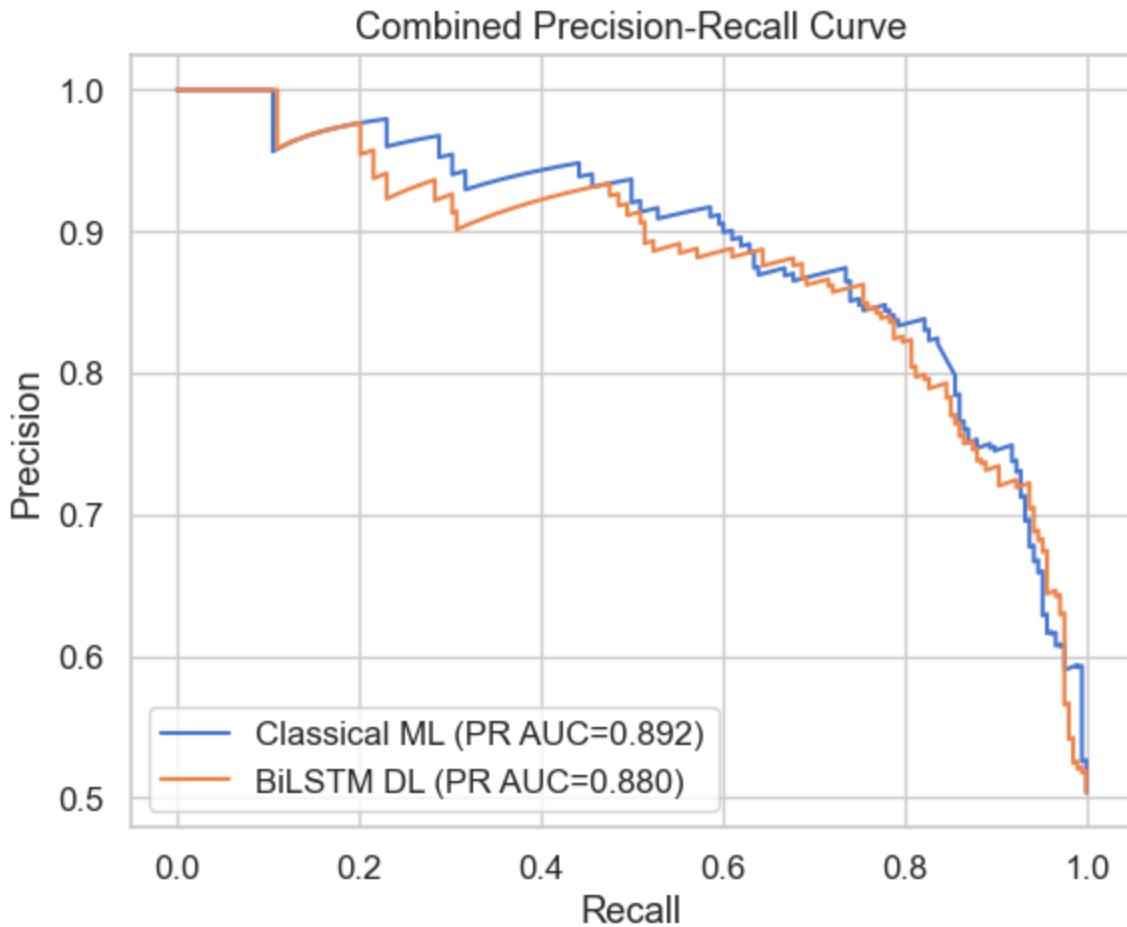


## 5.7 Combined Precision-Recall Curve

```
In [16]: plt.figure(figsize=(6,5))

# Plot Precision-Recall curve for Classical ML
plt.plot(recall_cls, precision_cls, label=f"Classical ML (PR AUC={pr_auc_cls:.3f})")
# Plot Precision-Recall curve for BiLSTM DL
plt.plot(recall_dl_vals, precision_dl_vals, label=f"BiLSTM DL (PR AUC={pr_auc_dl:.3f})")

# Labels and title
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("Combined Precision-Recall Curve")
plt.grid(True)
plt.legend()
plt.tight_layout()
plt.show()
```



## 5.8 Final Summary

```
In [17]: # Print a clear, final comparison table for both models
print("\n==== Final Model Comparison Summary ===")
print(df_metrics.to_string(index=False))

# Determine the best model based on F1-score (balances precision & recall)
best_model_name = df_metrics.loc[df_metrics['F1-score'].idxmax(), 'Model']
print(f"\nBest overall model based on F1-score: **{best_model_name}**")

==== Final Model Comparison Summary ====
    Model  Accuracy   Precision   Recall   F1-score   ROC-AUC   PR-AUC
Classical ML  0.808717  0.784141  0.855769  0.818391  0.891229  0.891547
      BiLSTM DL  0.803874  0.792627  0.826923  0.809412  0.882411  0.879646

Best overall model based on F1-score: **Classical ML**
```